

# Cloud Foundry: Open Source Alternative

Murali Cheruvu  
Indiana University  
3209 E 10th St  
Bloomington, Indiana 47408  
mcheruvu@iu.edu

## 1 INTRODUCTION

The cloud computing brings needed agility, scalability, storage, processing, global reach and reliability to software solutions. There are three popular models of cloud offerings: Software as a Service (SaaS), Platform as a Service (PaaS), and Infrastructure as a Service (IaaS). SaaS typically uses the internet to facilitate applications to its users. Some of the popular SaaS Offerings are: Google Apps, Dropbox and Salesforce. IaaS provides all the infrastructure needs - scalable and automated computing resources like servers and networking. Rackspace, Amazon Web Services (AWS), Microsoft Azure, Google Compute Engine (GCE) are some of the IaaS offerings. PaaS provides all the needed infrastructure, services and tools so that developers can focus on cloud-native development with aspects like micro-services and exchanging messages between the domain services. Examples of PaaS are: Microsoft Azure, Heroku and OpenShift. Cloud Foundry (CF) is, an open source, Platform as a Service (PaaS) available through a variety of private and public cloud distributions, which is developed and operated by VMware and then transferred to Pivotal Software, a joint venture by EMC, VMware and General Electric.

## 2 CLOUD FOUNDRY OVERVIEW

As opposed to most of the PaaS offerings in the market, which are tied to proprietary implementations, Cloud Foundry is an open source software with flexibility to allow integrations with external systems. Cloud Foundry provides all the PaaS capabilities like the popular public and private PaaS Providers - built-in scalable infrastructure, middleware, and various tools for development, deployment and support. To enable infrastructure-agnostic architecture, Cloud Foundry focused on three main categories: Clouds, Frameworks and Services.

### 2.1 Category: Clouds

Public and private clouds have their advantages and disadvantages. While public clouds provide flexibility and faster deployments, private clouds offer operational efficiency and total control. Hybrid cloud approach gives best of both public and private cloud offerings - infrastructure scalability, deployment and monitoring tools, data locality, industry regulations, zero-changes to the existing applications or develop with cloud-native mindset - API Gateways in combination of Micro-Services. Cloud Foundry is an open PaaS with ability to extend and collaborate with other private and public cloud systems. It can co-exist with other PaaS and IaaS platforms, in a way, it can be hosted on top of popular cloud environments like AWS, Azure, Google Cloud and OpenStack.

### 2.2 Category: Frameworks

Most of the cloud environments are restricted to fewer frameworks and programming languages. Though there is a good coverage of runtime environments and programming languages offered by most of the public clouds including AWS and Azure, there are still some restrictions either in the way we will need to use them or in the pricing and licensing models. Cloud Foundry, at architecture level, is generic and intended to host any programming language. It is written in Ruby and currently supports popular programming languages like Spring, Java, Ruby and NodeJS.

### 2.3 Category: Services

Similarly most of the clouds have limited set of support for various types of data, messaging and other services, restricting companies and development teams by forcing them to use specific technologies. Cloud Foundry, out of the box, shipped with support to various services like relational (MySQL, PostgreSQL), NoSQL (MongoDB), Key-Value pair (Redis) databases and RabbitMQ message queuing system along with extensibility to allow third-party systems to be added later.




Figure 1: Cloud Foundry - Open PaaS [?] ]

## 3 KEY COMPONENTS

Cloud Foundry comes with lots of ready-made components to support all the key aspects of PaaS cloud computing in a scalable fashion. Figure 2 is a good representation of the major components of Cloud Foundry:

## Cloud Foundry Architecture

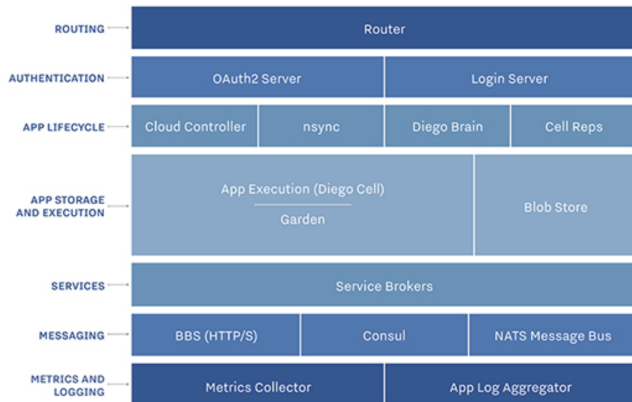


Figure 2: Cloud Foundry Architecture [? ]

### 3.1 Router

Router is responsible for controlling all the external and application level traffic, and also directing the incoming traffic to appropriate components. Router configuration allows to have number of routers to enable proper load balancing and high availability of the cloud foundry environment. Each router maintains a dynamic route table with all details of the deployed applications. Gorouter interacts with Cloud Controller and Droplet Execution Agent (DEA) to facilitate the updated routing information across the Cloud Foundry Environment. Router is implemented in Go programming language that can offer optimal performance.

### 3.2 UAA and Login Server

User Account and Authentication (UAA) and Login Server components are the identity management system in Cloud Foundry. Cloud Foundry uses OAuth2 (Open Authorization) standards driven token-based authentication and authorization to manage user security tokens.

### 3.3 Cloud Controller

We can think of Cloud Controller (CC) as the brain of Cloud Foundry Environment and the main responsibility of Cloud Controller is to manage the applications life cycle - deployment, application metadata, staging and running the applications. CC uses Diego Brain, CC-Bridge and Diego Cells to stage and run the applications. CC redirects the first requests to the appropriate Droplet Execution Engine (DEA) available in the load balancing pool. CC user permissions are maintained at various levels - Orgs, Spaces and Roles for greater scalability to the role-based access control. Application deployment artifacts - code packages, build-packs and droplets are maintained in Blob Store.

### 3.4 Execution and Storage

Droplet Execution Engine (DEA) is responsible for application deployment and runtime management - selecting appropriate build-pack, stage the application and ensuring end-to-end life cycle management of the application instance. Build packs are the scripts to identify the required framework for applications to run properly. Droplet is the unit of execution - deployed build pack of the application along with the application metadata. Wardens are containers to host the droplets and isolate them in resource-controlled environments. BOSH is used as tool for release management of complex distributed systems.

### 3.5 Service Brokers

Most of the applications have aspects like interacting with database, sending messages using Service-Oriented Architecture (SOA) and interfacing third-party components. Services cannot directly interact with the applications given they run in the containers which are not persistent. Cloud Foundry uses service brokers in a decoupled fashion through which application developers can facilitate and use the services in applications. CC uses NATS based messaging system. It uses a light-weight publish-subscribe mechanism to distribute the queued messages among applications.

### 3.6 Monitoring and Logging

CC provides various tools for continuous monitoring - (a) Health Manager: Responsible for monitoring health of the application instances by interacting with DEA and CC, (b) Metric Collector: Gathers various application metrics from the running instances, and (c) Log Aggregator: Streams application logs. Developers and support teams can access these logs and metrics to monitor, support and take necessary actions to keep systems up and running.

## 4 CONCLUSION

Lots of Fortune 500 companies from all over the world - various industries and government organizations, rely on Cloud Foundry for all the benefits of flexible open source PaaS offerings. Cloud Foundry is setting the new standards for cloud computing with emphasis on scalability and industry best practices by promoting 12-factor application development.

## ACKNOWLEDGMENTS

The author would like to thank Dr. Gregor von Laszewski and the Teaching Assistants for their support and valuable suggestions.

## REFERENCES